

REBOUND: Efficient, expressive, and well-scoped binding

Noé De Santo

ndesanto@seas.upenn.edu
University of Pennsylvania
Philadelphia, PA, USA

Stephanie Weirich

sweirich@seas.upenn.edu
University of Pennsylvania
Philadelphia, PA, USA

Abstract

We introduce the REBOUND library that supports well-scoped term representations in Haskell and automates the definition of substitution, alpha-equivalence, and other operations that work with binding structures. The key idea of our design is the use of first-class environments that map variables to expressions in some new scope. By statically tracking scopes, users of this library gain confidence that they have correctly maintained the subtle invariants that stem from using de Bruijn indices. Behind the scenes, REBOUND uses environments to optimize the application of substitutions, while providing explicit access to these data structures when desired. We demonstrate that this library is expressive by using it to implement a wide range of language features with sophisticated uses of binding and several different operations that use this abstract syntax. Our examples include `pi-forall`, a tutorial implementation of a type checker for a dependently-typed programming language. Finally, we benchmark REBOUND to understand its performance characteristics and find that it produces faster code than competing libraries.

Keywords: Dependent Haskell, well-scoped term representation, de Bruijn indices

1 Implementing binding

Implementors of programming languages, logics and calculi in Haskell have a choice to make when it comes to representing the binding structure of their programming languages. They need a representation of variables and binding locations (such as λ -expressions) that accurately represents their abstract syntax and operations that use this syntax (such as substitution, evaluation, and type checking). There are many binding representations possible [10, 13, 29, 35] and in this choice, implementors must balance multiple factors. In general, they would like one that is simple to work with, requires minimal boilerplate, and gives them confidence that their code is correct. At the same time, they would like an efficient implementation that does not slow down their code.

In the setting of mechanized programming language semantics, it is common to use de Bruijn indices with a *scope-safe* representation [2, 3, 6, 31, 33]. In this case, the abstract syntax tree uses a dependent-type index to statically track the number of free variables currently in scope. The types of

operations that work with syntax describe how they modify the current scope, and the type-checker statically verifies the correctness of this specification.

Haskell programmers have seen scope-safe representations of lambda calculus terms before, most notably in a functional pearl by Bird and Paterson [8], which uses a nested datatype to statically track scoping level, and in the Bound library [21], which optimizes this representation using an explicit weakening operation. However, despite the long history, this approach is not widely used in practice.

Therefore, we have developed the *REBOUND* library¹ as a tool to assist Haskell developers in working with well-scoped de Bruijn indices. This tool provides type classes, abstract data types, and can automatically derive necessary operations for working with variables. It is also accompanied by a suite of literate examples that demonstrate its use in various settings.

The design of this library is governed by three goals:

Correctness REBOUND uses Dependent Haskell to statically track the scopes of bound variables. Because variables are represented by de Bruijn indices, scopes are represented by natural numbers, bounding the indices that can be used. If the scope is 0, then the term must be closed. The type checker can identify when users violate the subtle invariants of working with indices.

Efficiency Behind the scenes, REBOUND uses first-class multisubstitutions, or *environments*, to delay the execution of operations such as shifting and substitution. Furthermore, these environments are accessible to library users who would like fine control over when these operations happen.

Convenience REBOUND is based on a type-directed approach to binding, where users indicate binding structure in abstract syntax through the use of types provided by the library. As a result, it provides a clean, uniform, and automatic interface to common operations such as substitution, alpha-equivalence, and free variable calculation.

Our goal with this paper is to highlight the key ideas that underlie the design of this library, to describe the design space and potential trade-offs in its implementation, and to evaluate its usability and performance at scale.

¹Available at <https://github.com/sweirich/rebound>.

In Section 2, we develop the key idea that underlies our approach: the use of parallel substitutions, that we call *environments* [1]. A well-scoped environment is a finite map from variable indices, bounded by n , to expressions with indices bounded by m . These two natural numbers are part of the environment’s static type and ensure that we only ever look up indices that are valid in the current scope, and that we know the scoping of the resulting term after the substitution has been applied. This section demonstrates that by considering how substitutions may be composed and delayed during the execution of an evaluator for the untyped lambda calculus, we are able to avoid multiple traversals over terms. Those choices can have a dramatic effect on the performance, as we demonstrate at the end of this section.

To make this idea readily available to Haskell programmers, in Section 3, we show that these key ideas can be packaged up inside appropriate type classes and abstract types, providing novice users with a simple interface to these operations. To evaluate the expressiveness of the REBOUND library and provide a tutorial on its usage, we have collected a suite of examples that challenge the capabilities of the library. We give an overview of this suite in Section 3.2 and demonstrate its support for *pattern binding*.

Furthermore, we have developed a case study to evaluate this work in the context of a practical setting. Section 4 discusses our adaptation of the implementation of the `pi-forall` language [36] to use this library. This code base includes a parser, scope checker and bidirectional type checker for a dependently-typed programming language with datatypes and dependent pattern matching.

When developing the library, we have choices about how we may *represent* environments in Haskell and with how functions that operate over lambda calculus terms may *use* environments as part of their operation. In Section 5, we deepen our performance analysis by benchmarking various environment representations and by comparing uses of REBOUND against other binding libraries available in Haskell. We also have benchmarked the performance of `pi-forall` using REBOUND against its prior implementation. In each case we find that REBOUND outperforms its competition, especially on benchmarks that require repeated β -reductions.²

The use of a well-scoped representation is form of dependently-typed programming in Haskell. While our examples and case studies provide positive evidence that this approach is beneficial, we acknowledge that there are trade-offs. We identify limitations with working with a scope-safe representation in Section 6. Finally, we discuss related work in Section 7 and conclude in Section 8.

²Our benchmarking code is also available at <https://github.com/sweirich/rebound>.

2 Well-scoped evaluators

Consider the following *scope-safe* representation of the lambda calculus:

```
data Nat = Z | S Nat -- Peano nats

data Fin :: Nat → Type where
  FZ :: Fin (S n) -- zero
  FS :: Fin n → Fin (S n) -- succ

data Exp :: Nat → Type where
  Var :: Fin n → Exp n -- variables
  Lam :: Bind n → Exp n -- abstractions
  App :: Exp n → Exp n → Exp n -- applications

data Bind n where
  Bind :: Exp (S n) → Bind n -- increase scope
```

This snippet first defines natural numbers which can be used at the type-level, and uses them to define bounded natural numbers (i.e. *finite naturals*), which will be used as de Bruijn indices [13] to represent variables. In terms, we use `f0`, `f1`, `f2`, etc. to refer to the (bounded) numbers 0, 1, 2, etc, when the bound can be inferred. The `Exp` type is indexed by a natural number, the *scope index*, so that it can only represent well-scoped expressions. All variables must be in the range specified by the scope of the datatype. The type of the `Lam` constructor states that the body of the expression is a binding, i.e it increases the scope by one. For reasons which will be explained later, we create a new type for binders for this increase. Note that, since `Fin 0` has no inhabitant, `Exp 0` cannot contain any free variable and hence represents *closed expressions*.

For example, to represent the closed lambda calculus term $\lambda x. \lambda y. y$, we use index `f0` for the occurrence of `y` as it refers to the closest enclosing binder. On the other hand, for $\lambda x. \lambda y. x$, the index of `x` in the same scope is `f1`.

```
example1 :: Exp 0 -- \x → \y → y
example1 = Lam (Bind (Lam (Bind (Var f0))))

example1 :: Exp 0 -- \x → \y → x
example1 = Lam (Bind (Lam (Bind (Var f1))))
```

2.1 An environment-based interpreter

Consider the implementation of an environment-based evaluator for the well-scoped representation. An *environment*, or closing substitution, is a mapping from variable indices to values. In the pure lambda calculus, a value is a closure [23]—an environment paired with the body of an abstraction.

```
type Env n = Fin n → Val -- environment type

data Val where
  VLam :: Env n → Bind n → Val -- closure
```

Environments can be constructed much like length-indexed lists.

```
nil :: Env Z -- empty env
(.:) :: Val → Env n → Env (S n) -- extend w/ value
```

An environment-based evaluator uses an environment argument to remember the values of variables.

```
eval :: Env n → Exp n → Val
eval r (Var x) = r x
eval r (Lam b) = VLam r b
eval r (App a1 a2) = case eval r a1 of
  VLam s (Bind b) → eval (eval r a2 .: s) b
```

The interpretation of a lambda expression is a closure. This value stores the current environment along with the body of the lambda expression. In the application case, this body is evaluated with the saved environment after it has been extended with the value of the argument of the application. Note that Haskell’s non-strict semantics gives this interpreter a call-by-need evaluation behavior—the argument is only evaluated if the variable is used in body of the abstraction.

There are two observations to make about this implementation. First, scope-safety means that the evaluator will never trigger a run-time error from an unbound variable. The environment type `Env` uses the scope index to statically track its domain, ensuring that every variable lookup is in scope. Second, there is no administrative work during evaluation. Even though we are using indices to represent variables, there is no shifting or substitution required. Instead, everything is handled via the environment.

However, despite these benefits there are also issues with this evaluator.

1. In our evaluation function, we need to explicitly pass around an environment and store a delayed environment in a closure. That doesn’t look like the lambda-calculus! What if we wanted something that looks more like a research paper, which often use substitution?
2. What if we wanted full normalization (i.e. reduction under abstractions) instead of evaluation (i.e. only top-level reduction of closed terms)?
3. The result of evaluation is a closure, not an expression. If we want to access the lambda calculus term corresponding to the closure we need to do more work. Furthermore, when comparing the results of evaluation, we do not wish to distinguish closures that differ in their saved environments.

The fundamental mechanism of this code is the *closure*, i.e. an expression that is paired with its environment. This environment acts as a *delayed* substitution, for significant benefit to our implementation. Let’s take this idea to see if we can bring some of the benefits of an evaluation-based interpreter to a substitution-based implementation.

2.2 A substitution-based interpreter

```
-- multisubstitution type
type Env m n = Fin m → Exp n
-- empty and "cons"
nil :: Env Z n
(.:) :: Exp n → Env m n → Env (S m) n
-- identity and composition
id :: Env n n
(.>>) :: Env m n → Env n p → Env m p
-- increment all variables by one
shift :: Env n (S n)
-- lift to a larger scope
up :: Env m n → Env (S m) (S n)
up e = Var f0 .: e .>> shift
```

Figure 1. Multisubstitutions and operations

But first, let’s consider a standard substitution-based implementation of an interpreter for the pure lambda calculus.

```
eval :: Exp n → Exp n
eval (Var x) = Var x
eval (Lam b) = Lam b
eval (App f a) = case eval f of
  Lam b → eval (instantiate b (eval a))
```

In this case, we don’t use an auxiliary type of values. Instead, evaluation, if it terminates, produces a new expression. The important step is in the application case: after evaluating the function, we substitute the evaluated argument into the body of the lambda term before its evaluation, using the function

```
instantiate :: Bind n → Exp n → Exp n
```

The instantiation function is defined through *substitution*, and a common implementation is shown in Figure 2a. The definitions in this figure rely on a small library (Figure 1) for working with mappings from indices to expressions, also known as *multisubstitutions*.

The use of *multi*-substitutions means that when applying a substitution to a lambda expression, it can simultaneously shift any free variables in its range to the new scope. To do so, we rely on the type for multisubstitutions, which we dub *environments* because they map indices in a bounded range. In contrast to the previous `Env` type, here the type cares about two numbers: m , the scope of the environment (i.e. size of its domain) and n , the scope of the expressions in the range. This is purely a type change; the environment extension operator is the same as above, but has a new type.

The `applyE` function applies the substitution to an expression. In the case of a lambda expression, the substitution must be *shifted* to work in the increased scope, via `up`. This operation modifies the substitution so that it leaves the bound variable alone (index `f0` is mapped to `Var f0`), offsets the rest of the substitution by one, and increments any free variables in the range of the substitution as they enter a new scope.

```

-- binder type
data Bind n where
  Bind :: Exp (S n) → Bind n

-- apply a multisubstitution to a binder
applyBind :: Env m n → Bind m → Bind n
applyBind r (Bind b) = Bind (applyE (up r) b)

-- apply a multisubstitution to a term
applyE :: Env m n → Exp m → Exp n
applyE r (Var x)   = r x
applyE r (Lam b)   = Lam (applyBind r b)
applyE r (App f a) = App (applyE r f) (applyE r a)

-- single substitution
instantiate :: Bind n → Exp n → Exp n
instantiate (Bind b) a = applyE (a .: Var) b

```

(a) Eager substitution

```

-- binder type with delayed substitution
data Bind n where
  Bind :: Env m n → Exp (S m) → Bind n

-- apply a multisubstitution to a binder
applyBind :: Env m n → Bind m → Bind n
applyBind r (Bind b) = Bind (r' .>> r) b

-- apply a multisubstitution to a term
applyE :: Env m n → Exp m → Exp n
applyE r (Var x)   = r x
applyE r (Lam b)   = Lam (applyBind r b)
applyE r (App f a) = App (applyE r f) (applyE r a)

-- single substitution
instantiate :: Bind n → Exp n → Exp n
instantiate (Bind r b) a = applyE (a .: r) b

```

(b) Delayed substitution

Figure 2. Eager and delayed substitutions

In contrast to the environment-based interpreter, working with substitutions requires bookkeeping. This bookkeeping costs in terms of performance (both substitution and shifting traverse the terms) and in terms of development time (the code in Figure 2a is slightly longer than the one in Section 2.1).

However, as above, this definition is scope-safe. The type of the substitution function tells us that it reduces the number of free variables in the term. The type of the evaluator restricts it to working with closed expressions.³

2.3 A delayed substitution-based interpreter

Now let's improve our substitution-based interpreter by using ideas from the environment based approach. The key technique is that of a *delayed substitution*, analogous to the closure above. Instead of eagerly substituting through the term, the term itself may contain an unapplied substitution.

One option would be to add an explicit substitution form to the expression datatype, following the $\lambda\sigma$ -calculus of Abadi et al. [1]. However, we can be a bit more sneaky, as the part of the term where this only really matters is at binders.

We modify the abstract type `Bind`, as shown in Figure 2b, so that it also contains a delayed environment. Note that, because we already specified the scope increase with the type `Bind`, the definition of `Exp` is not changed in any way. Thus, we are smuggling an environment into our expression type, hiding it behind an abstract type so that it does not need to be manipulated explicitly. This version of the `Bind`

type generalizes both lambda expressions and closures. If the delayed environment is `id`, which maps indices to corresponding variables, then this type is like a normal lambda abstraction. On the other hand, the type `Bind 0` is like the `Val` type from above and forces the delayed environment to be a closing substitution.

With these modifications, the implementation of the evaluator is identical to the version shown in Section 2.2.

```

eval :: Exp n → Exp n      -- same as in 2.2
eval (Var x)   = Var x
eval (Lam b)   = Lam b
eval (App f a) = case eval f of
  Lam b → eval (instantiate b (eval a))

```

The place where the delayed substitution comes into play is in the `applyBind` operation (Figure 2b). There, instead of shifting and applying the substitution to the body of the binder, we can wait by *composing* it with the suspended substitution in the binder, using the `(.>>)` operator. This observation was already present, in a slightly different form, in Bird and Paterson's functional pearl [8].

There is of course, no free lunch. By introducing the `Bind` type, we no longer have a unique representation for α -equivalent lambda expressions as they may differ in the substitutions suspended at binders. We account for this by equating the bodies of binders only after forcing their delayed substitutions.

³However, unlike before, the evaluator cannot statically guarantee that the result of evaluation (if any) will be a lambda expression, so there is the possibility of pattern match failure.

2.4 Explicit environments

We can delay the substitution even further by explicitly passing it as an argument to the evaluator, similar to the environment passing evaluator. This implementation fuses the traversal of the term during instantiation with the traversal of the evaluator itself.

```
eval :: Env m n → Exp m → Exp n
eval r (Var x)   = r x
eval r (Lam b)   = Lam (applyBind r b)
eval r (App f a) = case eval r f of
    Lam (Bind r' b) → eval (eval r a .: r') b
```

Compared to the previous definition, this version delays substitution even more, and ultimately does less work. With the previous version, in an application, we evaluate the body of the binder after substituting its (evaluated) argument for its parameter. That means that the (evaluated) argument gets re-evaluated again for each occurrence of the variable in the original body. Re-evaluation is fast, as this argument is already a value, but the revised version avoids it entirely. (While prior work has noted the benefit of delayed substitutions, we believe this version is new.)

2.5 What is the point of all of this?

In this section, we have considered four different evaluators (called EvalV, SubstV, BindV and EnvV respectively). These implementations are straightforward and directly map to our understanding of syntactic manipulations of the lambda calculus. However, these four evaluators perform differently, as shown by the table below. In each case, we timed the evaluation of the same large expression.⁴

Name	Detailed in	Description	Time
EvalV	Section 2.1	Original	315.7 μ s
SubstV	Section 2.2	Standard subst	3.546 s
BindV	Section 2.3	Delayed subst	816.3 μ s
EnvV	Section 2.4	Env argument	632.0 μ s
ExpSubstV	(not shown)	Explicit subst	6.255 ms

The first line of the table, the pure environment-based evaluator, is our baseline and produces the fastest time. The version with direct implementation of substitution (on the second line) is orders of magnitude slower. Notably, we can recover much of this time by merely delaying the substitution in the Bind term. Passing the environment explicitly brings us to about twice as slow as the original version. (But, note that these two versions also work with open terms, unlike

⁴The large expression was developed by Lennart Augustsson [4] and is the Scott encoding of `fact 6 == sum [0.. 37] + 17`. The benchmarks were run on a 2024 MacBook Pro M4 with 48 GB memory and timed using the criterion library. To observe the result of evaluation, we included boolean values `true` and `false` in language and extended the evaluators accordingly. These benchmarks are available in the directory `benchmark/lib/Rebound/Manual/Lazy`.

```
-- delayed substitution (abstract type)
type Env v m n
-- access environment at index m
(!) :: Env v m n → Fin m → v n
-- operations from Figure 1:
nil, (.), id, (.>), shift, up
-- identify the variable constructor
class Subst v v ⇒ SubstVar v where
    var :: Fin n → v n
-- apply environment to a term
class SubstVar v ⇒ Subst v e where
    applyE :: Env v m n → e m → e n
-- bind var v in body e (abstract type)
data Bind v e n
-- 'Subst' instance for 'Bind' (i.e. applyBind)
instance SubstVar v ⇒ Subst v (Bind v e)
-- single substitution
instantiate :: Bind n → Exp n → Exp n
```

Figure 3. Core library interface

```
-- identify the variable constructor
instance SubstVar Exp where var = Var

-- direct implementation of substitution
instance Subst Exp Exp where
    applyE r (Var x)   = r ! x
    applyE r (Lam b)   = Lam (applyE r b)
    applyE r (App f a) = App (applyE r f) (applyE r a)

-- implementation with generic programming
-- only need to identify the variable case
instance Subst Exp Exp where
    isVar (Var x) = Just (Ref1, x)
    isVar _      = Nothing
```

Figure 4. User code for well-scoped terms

the original evaluator.) We also compared these implementations with a fifth version, ExpSubstV, that more closely resembles the $\lambda\sigma$ -calculus [1] and allows suspended substitutions anywhere in the abstract syntax. However, on this benchmark that implementation is about ten times slower than EnvV.

3 The REBOUND library

We don't need to start from scratch in our next implementation of a language with binding. In this section, we separate the mechanism from the previous section into parts that are specific to the untyped lambda calculus and parts that can be reused for other languages and purposes and package that up in the REBOUND library.

Figure 3 isolates the library definitions necessary to implement the evaluation functions from Section 2. Then, Figure 4 uses these operations to implement substitution for the untyped lambda calculus twice: first directly and then by deriving the definition using generic programming.

This first `Subst` instance is simple because the library already includes an instance for applying the substitution to a binder: the composition and delaying of the `Env` type happens behind the scenes. The second instance only requires the user to identify the variable case in the abstract syntax (if there is one), but requires no modification when new syntactic forms are added to the language.

REBOUND keeps the `Env` type abstract. While one way to implement delayed substitutions is with functions, as shown in the previous section, that is not the only possible implementation. We discuss alternatives in Section 5.2. Because this type is abstract, we include an explicit operator `!` for looking up an index in the environment.

The `SubstVar` class identifies scope-indexed types that have variable constructors. The `Subst` type class takes two arguments. The first, `v`, describes the co-domain of the deferred substitution (i.e. what type do variables stand for) and the second `e` describes the type we are substituting into. Often, these two types will be the same, e.g., in Figure 4, we instantiate both parameters of the `Subst` class with `Exp`.

The abstract `Bind` type. The type of single binders (`Bind`) is abstract and the library includes relevant type class instances for this type, such as `Subst`. Internally, the `Bind` type includes a suspended environment, as in Figure 2b, but users need not be aware of this delayed substitution. Instead, they should work with the `bind` and `getBody` wrappers.

```
bind :: SubstVar v => e (S n) -> Bind v e n
bind = Bind id
```

```
getBody :: Subst v e => Bind v e n -> e (S n)
getBody (Bind r e) = applyE (up r) e
```

REBOUND also includes operations that allow users to manipulate environments explicitly. For example, a user may wish to instantiate a binder while calling a function that is parameterized by the current environment.

```
instantiateWith :: (forall m n. Env v m n -> e m -> e n)
               -> Bind v e n -> v n -> e n
instantiateWith f (Bind r a) v = f (v .: r) a
```

This library function is exactly what is required to implement the environment-based interpreter shown in Section 2.1 while keeping the `Bind` and `Env` types abstract.

3.1 Beyond the untyped lambda calculus

Many languages include rich binding structures. We would also like to implement more functions than evaluators, such as normalizers (which reduce open terms) and type checkers.

Finally, we would like to use this library in full-featured implementations, so it must be compatible with their additional requirements.

To demonstrate the features of this library, we have used it to represent the binding structure for a number of different calculi, and have implemented normalizers and type checkers for these languages. These examples have been extensively documented and are distributed along with the library.

LC.hs Untyped lambda calculus with single binding. Big-step and small-step evaluation functions using substitution, normalization.

LClet.hs Untyped lambda calculus with `let` binding, which may be recursive or nested. Big-step evaluation and normalization.

Pat.hs Untyped lambda calculus with constants and pattern matching. Big-step and small-step evaluation.

SystemF.hs System F with separate term and type variables. Type checker.

PureSystemF.hs System F with a unique syntactic class for terms and types. Type checker.

PTS.hs Dependently-typed calculus including Π and weak Σ types, based on Pure Type Systems [5]. Big-step and small-step evaluation, normalization. Bidirectional type checker.

DepMatch.hs Dependently-typed calculus with nested dependent pattern matching for strong Σ types. Big-step and small-step evaluation, normalization. Bidirectional type checker.

We also have a few examples that demonstrate how to work with well-scoped expressions.

ScopeCheck.hs Scope checker: converts a nominal representation of binding to a well-scoped version.

LCQC.hs Generator for well-scoped untyped lambda calculus terms, suitable for property-based testing using the QuickCheck [12] library.

HOAS.hs Uses HOAS as a convenient interface to construct concrete well-scoped expressions.⁵

PatGen.hs A version of `Pat.hs` that demonstrates the use of generic programming in the presence of sophisticated scoped-indexed types.

3.2 Pattern binding

Single binders work well for theoretical developments. But we often want more from a binding library in a practical implementation, such as *pattern binding*. A pattern can be any type: all we need to know about it is how many variables it binds. Figure 5 shows the generalized interface for the `Bind` type from a single index to *pattern binding*.

The type class `Sized` describes types that statically identify the number of variables that they bind, using the associated type `Size`. This class also includes the function `size` that returns the same information as a *singleton type* [15]. The type

⁵This example is inspired by McBride’s classy hack: <https://mazzo.li/epilogue/index.html%3Fp=773.html>

```

661 class Sized (t :: Type) where
662   -- retrieve size from the type (number of variables
663   -- bound by the pattern)
664   type Size t :: Nat
665   -- access size as a natural number term
666   size :: t → SNat (Size t)
667
668   -- bind variables for v, in expressions c
669   -- with patterns p in scope n
670   type Bind v c (p :: Type) (n :: Nat)
671
672   -- create a binder for the pattern p, introducing
673   -- its variables into the scope
674   bind :: (Sized p, Subst v c) ⇒ p → c (Size p + n)
675         → Bind v c p n
676
677   -- instantiate a binder by filling in values for
678   -- the variables bound by the pattern
679   instantiate :: (Sized p, Subst v c) ⇒ Bind v c p n
680               → Env v (Size pat) n → c n

```

Figure 5. Pattern binding interface

SNat n contains a natural number isomorphic to n . Because we lack true dependent types in Haskell, singleton types provide a bridge between runtime and compile-time data.

n -ary binding. The simplest form of pattern binding, is binding several variables at once. For example, the language of PTS.hs eliminates products using pattern matching instead of projections. It includes a “split” term that simultaneously binds *two* variables to the two components of the pattern. Therefore, it uses the singleton type SNat 2 as a pattern that binds exactly two variables.

```

696 data Exp n where
697   ... -- other constructors as before
698   -- create a product `(e1, e2)`
699   Pair :: Exp n → Exp n → Exp n
700   -- split a product `let (x,y) = e1 in e2`
701   -- the body of the binder has extended scope (2 + n)
702   Split :: Exp n → Bind Exp Exp (SNat 2) n → Exp n

```

Because the number of bound variables can be static determined from the pattern, the Bind constructor in Figure 5 increases the scope of the body of the binder by the number of variables bound in the pattern, and requires the same number of values in instantiation. Continuing the example, we extend the evaluator with cases for Pair and Split as below. In the latter case, the type checker requires us to supply two arguments to instantiate, packaged in an environment.

```

712 eval (Pair a1 a2) = Pair a1 a2
713 eval (Split a b) = case eval a of
714   Pair a1 a2 →

```

```

716 eval (instantiate b (a1' .: a2' .: nil))
717   where a1' = eval a1
718         a2' = eval a2

```

Nested pattern matching. Pattern binding also extends to arbitrary datatype patterns and nested pattern matching. For example, suppose we would like to add the ability to deeply match tuples in let bindings, i.e. `let (x, (y, z)) = e1 in e2`. To do so, we can define a datatype to represent the tuple structure of the pattern (Pat) and use this pattern in a new expression form (LetPair).

```

727 data Pat (m :: Nat) where
728   PVar :: Pat N1 -- binds a single variable
729   PPair :: Pat m1 → Pat m2 → Pat (m2 + m1)
730
731 data Exp (n :: Nat) where
732   ... -- other constructors as before
733   LetPair :: Exp n → Bind Exp Exp (Pat m) n → Exp n

```

Above, a pair pattern is either a single variable or an application of the pair constructor to two nested patterns. To know how many arguments are bound in this case, we sum the number of binding variables in each subpattern.

The evaluator for LetPair expressions must first identify whether the pattern matches a given value, and if so, produce a substitution for each of the variables in the pattern to the corresponding subterms in the value.⁶

```

743 eval (LetPair a b) = eval (instantiate (getBody b) r)
744   where
745     r = patternMatch (getPat b) (eval a)
746
747 patternMatch :: Pat p → Exp m → Env Exp p m
748 patternMatch PVar e = e .: nil
749 patternMatch (PPair p1 p2) (Pair e1 e2) =
750   withSNat (size p2) (r2 .++ r1) where
751     r1 = patternMatch p1 e1
752     r2 = patternMatch p2 e2

```

For PPair, we use an environment append operation, `(.++)`, to combine the results of pattern matching the components of the pair. This operation implicitly needs the length of its first argument at runtime; the function withSNat uses a value of type SNat n to satisfy this constraint.

This idea can also be used to implement pattern matching for arbitrary datatypes, as we demonstrate in the example Pat.hs.

4 Case study: pi-forall

To test the expressiveness of our library, we have ported pi-forall [36], a demo implementation of a type checker for a dependently typed programming language, to REBOUND. The pi-forall implementation includes a parser and type

⁶For simplicity, the code above throws an error when the expression has the wrong form; a more realistic example would gracefully handle this case.

checker for a language with dependent functions, datatypes, dependent pattern matching, multiple modules and informative error messages. The original implementation used the unbound-generics [20, 37] library (called unbound for short) to implement substitution and alpha-equivalence. This binding library relies on a *locally nameless representation*.

The previous implementation of `pi-forall` did not statically track the scoping of variables, relying instead on unbound’s design to ensure a correct treatment of binders. Therefore, we were curious to learn whether REBOUND, and more generally intrinsically scoped representations, could be used in a setting that is closer to a practical implementation. While `pi-forall` is a tutorial, focusing more on explaining how dependent types work than on developing a robust and efficient language, the features of `pi-forall` make it more than just a toy example.

The goal of this re-implementation is to evaluate the expressiveness of the core library by using it to implement a non-trivial programming language. As part of this process, it provided practical motivation for new extensions of the library. In particular, two features of `pi-forall` provide the greatest challenge to REBOUND; we discuss them below.

Error messages: Feedback from the type-checker to the user is crucial in dependently-typed languages as the types become expressive/complex. This feedback comes in the form of type errors, warnings, and a special `PRINTME` expression that instructs the type-checker to print the types of all variables currently in scope. It is important for such feedback to refer to variables as they were defined by the user and not using their index in the (current) scope. This means that `pi-forall`’s type-checker must maintain a mapping from de Bruijn indices to user-defined names during all parts of the implementation.

Datatypes and pattern matching: A significant amount of code in the `pi-forall` type checker involves checking datatype declarations, uses of type constructors and data constructors, and pattern matching expressions. Supporting `pi-forall`’s indexed datatypes requires expressive support for *telescopes*, sequences of variable declarations where the type of each identifier may refer to any variable bound earlier in the telescope. Telescopes are complex patterns, as they bind variables both internally (later in the same telescope) and externally (in the subsequent expression). Scope safety means that two values must be tracked: the number of variables bound by the telescope, and the (extending) scope for its embedded expressions.

4.1 Scoped Monads

When working with abstract syntax, Haskell programmers often use a `ReaderMonad` [19] to store information about in-scope variables, such as user-supplied names, types, or definitions. However, when working with de Bruijn indexed terms and statically tracking the current scope, the usual

reader monad is not sufficiently expressive. For example, when storing a length-indexed vector of the names of variables currently in scope, we might like to define an operation for extending that scope (i.e. consing a new name to the vector).

```
-- a simple monad that tracks names currently in scope.
-- Note that the scope is part of the type!
type Scoped n = Reader (Vec String n)
```

```
-- a specialized version of local
addToContext :: String → Scoped (S n) a → Scoped n a
addToContext x = local (x :>) -- Type error!
```

The expression `(x :>)` has type

```
Vec String n → Vec String (S n)
```

but `local`’s type requires a function with type

```
Vec String n → Vec String n
```

Therefore, to maintain this information, REBOUND defines the `ScopedReaderMonad` class. As its name implies, this monad offers the same API which allows to read and update a piece of data, usually called the monad’s *environment*. The key difference is in the types: the `ScopedReaderMonad`’s environment has to be indexed by a scope, and the `localS` operation is allowed to change the scope.

```
class (∀ n. Monad (m n))
  ⇒ MonadScopedReader (e :: Nat → Type) m | m → e
where
  -- | Retrieves the monad environment.
  askS :: m n (e n)
  -- | Executes a computation in a modified environment.
  localS :: (e n → e n') → m n' a → m n a
  -- | Retrieves a function of the current environment.
  readerS :: (e n → a) → m n a
```

By defining the `Scoped` monad so that it is an instance of this class, the `localS` method has the type that we need.

```
instance MonadScopedReader (Vec String) Scoped
  where ...
```

```
addToContext :: String → Scoped (S n) a → Scoped n a
addToContext x = localS (x :>) -- Type checks!
```

4.2 Scoped patterns and Telescopes

The most significant issue with datatype definitions in `pi-forall` is that the telescopes for constructors both bind new variables and include occurrences of existing variables. For example, the usual length-indexed vector can be expressed in `pi-forall` using the following top-level declaration.

```
data Vec (A : Type) (n : Nat) : Type where
  Nil   of [n = Zero]
  Cons  of [m : Nat] (h : A) (t : Vec A m) [n = Succ m]
```


This declaration includes a telescope for the parameters of the `Vec` type (i.e. `A` and `n`) and a telescope (in the scope of the first one!) for the parameters of each constructor (e.g. `m`, `h`, and `t` for `Cons`). The telescopes for constructors may also include constraints (or “Ford equations” [28]) on the parameters, such as `n = Zero` in the `Nil` case, constraining the length to be `Zero` for empty vectors.

This dual treatment of variables means that the simple datatypes for patterns, presented in Section 3.2, are not expressive enough. Instead, we need to statically track both the number of bound variables and the current scope. In other words, we use patterns of kind $\text{Nat} \rightarrow \text{Nat} \rightarrow \text{Type}$ instead of $\text{Nat} \rightarrow \text{Type}$, where the first argument is the number of bound variables and the second argument is the current scope regulating free variables.

Furthermore, in a binding telescope, variables bound earlier in the telescope can occur in types and constraints that appear later in the telescope. Re-using `Cons` as an example, its telescope binds the variable `m` and then uses it as both the length of the sublist `t`, and in the constraint on `n`.

The `REBOUND` library defines the `TeleList` datatype to support telescopes. A `TeleList` is parameterized by both `p`, the number of variables that it binds and `n` the scope that it appears in. It is also generic over `pat` a similarly parameterized pattern type for each entry in the telescope. In the `TNil` case, the telescope binds no variables `N0` and is available in any scope (`n` is unconstrained). However, in the `TCons` case, if the entry binds `p1` variables, then the rest of the telescope occurs in the extended scope `p1 + n`. Furthermore, the number of variables bound by the telescope includes both those bound here in the head and those bound later in the tail.

```
data TeleList (pat :: Nat → Nat → Type) p n where
  TNil  :: ( ... ) ⇒ TeleList pat N0 n
  TCons :: ( ... ) ⇒
    pat p1 n → TeleList pat p2 (p1 + n) →
    TeleList pat (p2 + p1) n
```

Using this generic definition, `pi-forall`’s telescopes can be defined by first a `Local` type describing an element of the telescope, and then applying `TeleList` to it:

```
data Local p n where
  -- Variable binding, e.g. (h: A)
  LocalDecl :: LocalName → Typ n → Local N1 n
  -- "Ford" constraint, e.g. [n = Succ m]
  LocalDef  :: Fin n → Term n → Local N0 n
```

```
type Telescope = TeleList Local
```

This type represents either a local variable declaration or an equality constraint. In the former, the pattern binds one variable, and the type of that variable is in scope `n`. In the latter, no variables are bound, but the equation must have a variable in the current scope on the left-hand side, and a term in the current scope on the right hand side.

Table 1. Comparison of normalization benchmarks

Benchmark Name	eval	nf	random15
Env.Strict.BindV	1.179 ms	1.384 ms	696.6 μ s
Env.Strict.EnvV	639.1 μs	972.0 μs	585.2 μs
Env.Strict.EnvGenV	831.2 μ s	1.300 ms	767.6 μ s
Env.Strict.Bind	5.005 ms	5.169 ms	676.8 μ s
Env.Strict.Env	669.0 μ s	1.000 ms	595.3 μ s
Env.Strict.EnvGen	866.3 μ s	1.347 ms	803.8 μ s
DeBruijn.BoundV	1.134 ms	1.236 ms	3.995 ms
DeBruijn.Bound	4.335 ms	4.389 ms	3.778 ms
Named.Foil	172.2 ms	173.1 ms	208.6 ms
Unbound.Gen	1.919 s	1.898 s	17.38 ms
Unbound.NonGen	1.199 s	1.177 s	3.157 ms
NBE.KovacsScoped	347.9 μ s	346.2 μ s	89.12 μ s

The constructors also includes constraints (elided) that help Haskell’s type checker work with telescopes. For `TNil`, the constraint states that $n+0 = n$. The `TCons` constructor has two. The first states that the size of the pattern is independent of the scope in which it appears. More formally,

$$\forall n. \text{Size}(\text{pat } p \ n) \sim p$$

The second asserts an associativity property about addition, instantiated with the binding variables. The `REBOUND` library includes smart constructors to supply these two constraints automatically. These facts are both brought into scope whenever the telescope is pattern matched, so are automatically available to the type checker during traversal of the telescope.

5 Benchmarks

Here we justify our claim that `REBOUND` provides an efficient implementation of interpreters and type checkers. As we report in this section, we have developed two sorts of benchmarks: *normalization* and *dependent type checking*.

The normalization benchmarks are a broad comparison across multiple implementations of lambda calculus normalization. We use these benchmarks to compare different ways of using `REBOUND`, different libraries for binding (Section 5.1), and different implementations of `REBOUND`’s environment data structure (Section 5.2). The dependent type checking benchmarks (Section 5.3) compare the performance between two versions of `pi-forall`. They model a more realistic language and draw on several different operations on syntax working together in a more realistic usage.

5.1 Normalization benchmarks

This section compares implementations of normalization for the untyped lambda calculus expressions. Our normalization function, `nf` fully reduces its argument, including underneath binders. It is defined in terms of an auxiliary function `whnf`,

that calculates the *weak-head normal form* of an expression, i.e. reduces just enough to reveal the top-level structure.

Table 1 shows the results on various tasks.

eval Weak-head reduction of Augustsson’s term encoded in the untyped lambda calculus extended with boolean values. This is the same benchmark used in Section 2.5.

nf Full normalization of Augustsson’s original term.

random15 Full normalization of a collection of 100 randomly generated terms that need *at least* 15 steps to normalize.

REBOUND implementations. The first six lines of Table 1 are implementations of full reduction using REBOUND. These include `Env.Strict.Bind`, the analogue to the delayed substitution implementation of Section 2.3, and `Env.Strict.Env`, the analogue to the explicit environment implementation of Section 2.4. For this section, we do not include analogous of `Env.Lazy.EvalV`, `Env.Lazy.SubstV` or `Env.Lazy.ExpSubstV`. We omit the first because we would like to measure full normalization, which is not supported by that interpreter. The latter two are omitted because they are orders of magnitude slower during evaluation.

In this benchmark set, it makes a (small but measurable) difference whether the abstract syntax trees used to represent lambda calculus terms are strict or non-strict. In Section 2, we used a non-strict representation for simplicity. Here, for uniform comparison, we exclusively use strict abstract syntax (both for REBOUND and other implementations).

We also explored the impact of reducing the argument before beta-reduction. In other words, whether we instantiate in both the `nf` and `whnf` functions with `a` or with `whnf a`. The names of benchmarks that use `whnf a` end with `V`. In `nf`, this modification is practical only when `a` is evaluated lazily; as normalizing all subexpressions can cause significant blow-up.

Finally, we modified the explicit environment versions (resulting in `Env.Strict.EnvGenV` and `Env.Strict.EnvGen`) to use GHC.Generics so that we can measure the cost of generic programming. Roughly, we found a 30-25% cost for this convenience.

Other implementations. To compare how REBOUND stacks up, we compared its performance against several other libraries (Section 7 describes these libraries in more detail).

DeBruijn.Bound defines well-scoped de Bruijn indices using Edward Kmett’s Bound library [21].

Named.Foil uses a nominal representation of lambda calculus terms. This code was developed by the foil library authors [25].

Unbound.Gen uses generic programming via unbound [20, 37]. The version `Unbound.NonGen` defines relevant operations by hand.

The results appear in the bottom half of Table 1. Overall, the REBOUND-based implementation `Env.Strict.EnvV` is

Table 2. Environment implementations

Benchmark Name	eval	nf	random15
main/Functional	1.150 ms	126.6 ms	829.6 μ s
main/Lazy	641.9 μ s	973.8 μ s	581.8 μ s
main/LazyA	640.2 μs	970.8 μs	575.9 μs
main/LazyB	1.509 ms	116.4 ms	965.8 μ s
main/Strict	720.1 μ s	1.019 ms	621.6 μ s
main/StrictA	720.3 μ s	1.039 ms	635.7 μ s
main/StrictB	1.678 ms	157.0 ms	1.002 ms
nat-word/Lazy	677.4 μ s	1.136 ms	608.0 μ s
vector/Vector	2.972 ms	1.534 s	2.178 ms

the fastest. The Bound-based implementations are competitive for `eval` and `nf`, but significantly slower on the random terms. Because the Foil and unbound versions do not delay substitutions, they are significantly slower on all benchmarks.

The bottom of the table includes a normalization function developed by Kovács (and modified to use well-scoped expressions by the authors).⁷ This version does not use the same algorithm—instead it uses *normalization-by-evaluation*, an alternative that is not based on substitution, and so does not generalize to other operations on lambda-calculus terms such as type checking or compiler optimization. We include it as baseline comparison with a fast algorithm, and it is consistently the fastest version. This suggests that when performance is critical, programmers can still use well-scoped representations, but may wish to look for specialized algorithms for the normalization part of their code base.

5.2 Environment implementation

Figure 3 defines the interface for the REBOUND library using an abstract type for environments (`Env`). As part of our benchmarking, we compared `Env.Strict.EnvV` compiled with several different implementations for this data structure.

main/Functional is the simplest implementation and represents environments as functions of type $\text{Fin } n \rightarrow \text{Exp } m$. **main/Lazy** is the implementation that we use in the library (and for the benchmarks in Table 1). This implementation “defunctionalizes” the environment as a data structure, representing environment creation functions (`idE`, `(. :)` `shift`, etc) as constructors. As a result, operations such as composition can perform optimizations, such as those found in Abadi et al. [1]. Furthermore, applications of the identity substitution can be optimized away [34].

main/LazyA This version is the same as `main/Lazy`, but does not include the optimized identity application.

⁷<https://github.com/AndrasKovacs/elaboration-zoo/blob/master/01-eval-closures-debruijn/Main.hs>

main/LazyB This version is the same as `main/Lazy`, but does not include the optimized construction.

main/Strict,main/StrictA,main/StrictB These versions are the same as `main/Lazy` and its variants, but use a strict spine.

These benchmarks reveal that the defunctionalized and optimized version is faster than using a function to represent the environment, and that much of the speed up comes from “smart composition”. Indeed, in the lazy version, the optimized identity application slightly degrades performance. The strict variants are also slightly slower.

The next benchmark (**nat-word/Lazy**) replaces runtime natural numbers (`SNat` and `Fin`) with machine words in **main/Lazy**. Finally, **vector/Vector** represents environments using `Data.Sequence` in addition to using machine words.

Overall, and somewhat surprisingly, the performance for the version with machine words is slightly worse than its unary analogue, and the results for `Data.Sequence` are significantly worse than our original version. There could be multiple factors at play. The `vector/Vector` implementation cannot take advantage of the optimizations of the defunctionalized version. Furthermore, while this data structure remains lazy in the elements stored in the environment, the structure of the vector is computed strictly. As a result, representing a shifting environment in a large scope takes much more space. Finally, this version needs a runtime representation of the current scope consistently throughout, requiring additional overhead.

5.3 Type checking benchmarks

The next set of benchmarks compares the new implementation of `pi-forall`, described in Section 4, with the original implementation. These benchmarks take the form of short (at most 250 lines) `pi-forall` programs. The time shown in Table 3 is the end-to-end time to process the file.

AVL An AVL tree implementation. This implementation is standard and doesn’t rely on dependent-types.

DepAVL An AVL tree implementation that internally uses dependent-types to enforce the AVL invariants.

Compiler A compiler from intrinsically typed arithmetic expressions to an intrinsically typed stack language, plus interpreters for both languages.

Lennart An adaptation of the benchmark from Section 2.5.

CompCk Checks that directly interpreting a program computing the factorial of 8 and interpreting the compiled program yield the same result.

According to these benchmarks, the REBOUND-based implementation is the fastest. The gap is most stark on the “compute intensive” benchmarks, but the difference is noticeable across the board.

The last two benchmarks could be considered synthetic, as they are not programs one would typically write. However they demonstrate that heavy computations can occur in

dependently-typed languages at the type-level. Furthermore, it is reassuring to observe that the performance difference between REBOUND and unbound also occurs in more fleshed out setting.

6 Is scope-safety a win?

REBOUND uses a *scoped* representation of terms, where the current scope is tracked statically by the type system. As a result, the operations supported by this library have expressive types, which can help users avoid bugs.

However, there is a cost associated with working with a scope-safe representation in terms of development time and flexibility. While more expert users, familiar with dependently-typed programming in Haskell, may benefit from the enhanced static checking, novice Haskell programmers may struggle with the complexity of the interface. Furthermore, scope safety may not always be the most appropriate representation choice, due to the limitations that we list below.

Reasoning about natural numbers When working with well-scoped terms, we need to prove to the type checker that scopes line up, which involves reasoning about the equality of natural number expressions. Thus far, we have wanted to record exactly what properties are needed that do not follow directly from their definitions. Therefore, we have avoided the use of a special purpose solver; instead the library includes two monoid axioms about natural number addition, which must be used explicitly.

Type inference This work includes many operations that are polymorphic over types of kind `Nat → Type` (i.e. scope-indexed types). However, type inference is more challenging when working with scoped patterns. There we use associated types [9] to indicate the number of variables bound by the pattern because the parameter already refers to the scope of expressions that appear inside the pattern itself. Unfortunately, associated types interact poorly with unification and type class resolution. Our example suite includes several positive examples to demonstrate the appropriate annotations needed to guide the Haskell type checker.

Multiple binding sorts Our library is scope-safe for a *single* scope and does not support multiple sorts of scopes well. This causes difficulties for languages such as System F [17], that bind both type and term variables. While it is possible to index expressions by two different scopes, using the library requires conversions to make sure that the “right scope” is the last parameter at certain times. Alternatively, users can combine both type and term binding in the same scope. REBOUND includes examples of both approaches.

Generic programming REBOUND uses `GHC.Generics` to automatically derive substitution and other operations. However, scope-safety imposes two complications, which our example `PatGen.hs` demonstrates how to resolve. First, generic

Table 3. pi-forall implementation comparison

	AVL	DepAVL	Compiler	Lennart	CompCk
Unbound	27 ms	47.3 ms	27 ms	1.83 s	1.69 s
Rebound	21 ms	36 ms	22.8 ms	46.4 ms	184 ms

programming only available for datatypes that do not include any “existential” variables. By isolating existentials into separate small datatypes, users can provide these instances by hand while still retaining the benefits of generic programming for the rest of their data structure. Second, all type constructors used in the definition of the syntax must include their scope as the last argument. REBOUND provides an alternative definition the list type to represent sequences of scoped expressions a scoped type.

7 Related Work

There are several binding libraries available for Haskell. The most similar to this work is Bound [21], which represents lambda calculus terms using well-scoped de Bruijn indices [8]. In contrast to REBOUND, which makes extensive use of Dependent Haskell features, users of Bound represent syntax using a nested datatype and uses a (derived) monad instance for the type as the (single) substitution operation. Bound adds additional support for statically tracking scopes and optimizing the implementation by delayed shifting. Examples distributed with the library show that it can be used with debug names and in a language with pattern matching.

The unbound library [20, 37] uses a *locally nameless* representation for variable binding [29]. In this approach, bound variables are represented by de Bruijn indices and free variables are represented by names (i.e. strings). When entering a new scope, users must replace bound variables with fresh free variables. It also develops a library of pattern types to assist in the definition of pattern binding. It does not track the scopes of free variables statically. The modern implementation of the library [20] uses generic programming [26] to automate the definition of these operations and a freshness monad to generate free names.

The *Foil* library [25] uses phantom types to track the scopes when variables are represented using *names*. This use of names is based on *the rapier* [30], the approach used internally by the GHC compiler. The key invariant is that names must be unique within their scopes; on entering a new scope, binders must be renamed if they have already been used. The *Free Foil* [22] extension adds support for pattern matching and uses TemplateHaskell [32] to automate boilerplate definitions. To evaluate its expressiveness, foil’s authors have used it to implement lambda-pi, a tutorial dependently-typed language [24].

Several authors also describe how to implement binding structures in Haskell. Augustsson’s note [4] provides simple

Haskell implementations using names, de Bruijn indices and higher-order abstract syntax (HOAS). HOAS based embeddings can be modified using the type system to rule out exotic terms [27, 35], and track scopes and types. Bernardy and Pouillard [7] describe methodology for a scope-safe higher-order interface layered on top of a de Bruijn indexed representation.

It is also possible to work with a well-typed representation of syntax in Haskell, which extends scope-safety with additional typing constraints for the object language. Guillemette and Monnier demonstrates the encoding of a type-preserving compiler [18]. Eisenberg’s Stitch functional pearl [14] includes a parser, type checker and optimizer for a statically typed core language. The Crucible language [11] uses a well-typed core language to build a suite of static verifiers.

Going outside of Haskell, Cockx⁸ provides an overview of variable representations that are possible in the dependently-typed language Agda. While our approach based on well-scoped de Bruijn indices draws on similar work done in the context of a proof assistant [2, 3, 6, 16], there are two important differences. First, type theories such as Agda, Rocq and Lean, require showing that substitution functions are total. This must be done by decomposing it into two steps: first renaming and then substitution. In Haskell, we can define substitution in one go. Second, when environments are delayed binders, α -equivalent terms do not have a unique representation, depending on what environment is stored in the term. Therefore, additional care must be taken to make sure that all judgments are stable up to this equivalence.

8 Conclusion

The REBOUND library provides a framework for working with binding structures that is, we believe, approachable to Haskell programmers. By statically tracking scopes, REBOUND eliminates many sources of confusion when working with de Bruijn indices. Our library design isolates the complexity of binders using an abstract type and automates the development of syntactic operations through generic programming. This approach is expressive, as we have demonstrated the by using REBOUND to implement languages with many different binding structures, and have implemented several different operations using this representation. These

⁸<https://jesper.sikanda.be/posts/1001-syntax-representations.html>

well-documented examples are part of the REBOUND repository, and form an extensive tutorial on working with well-scoped representations using Dependent Haskell. Accompanying the examples is an extensive case study that demonstrates an end-to-end use of the library in a practical setting. Finally, we have evaluated the efficiency of REBOUND against competing approaches and have found that it outperforms its competitors, sometimes significantly. After this thorough evaluation, we see no impediment to using this library at scale, in an industrial-strength system.

There are a number of potential avenues for future work. In particular, we would like to develop and benchmark a more extensive set of examples that demonstrate the use of this library for working with syntax. We would also like to explore methodologies to improve error reporting when code fails to type check. Finally, we would like to explore the use of a type-checker plug-in to automate reasoning about natural number scopes.

Acknowledgments

This material is based upon work supported by the National Science Foundation under Grant No. No. XXXXX. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the National Science Foundation.

References

- [1] M. Abadi, L. Cardelli, P.-L. Curien, and J.-J. Lévy. 1991. Explicit substitutions. *Journal of Functional Programming* 1, 4 (1991), 375–416. doi:10.1017/S095679680000186
- [2] Guillaume Allais, Robert Atkey, James Chapman, Conor McBride, and James McKinna. 2021. A type- and scope-safe universe of syntaxes with binding: their semantics and proofs. *Journal of Functional Programming* 31 (2021), e22. doi:10.1017/S0956796820000076
- [3] Thorsten Altenkirch and Bernhard Reus. 1999. Monadic Presentations of Lambda Terms Using Generalized Inductive Types. In *Computer Science Logic*, Jörg Flum and Mario Rodríguez-Artalejo (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 453–468.
- [4] Lennart Augustsson. 2006. λ -calculus cooked four ways. (2006).
- [5] H. P. Barendregt. 1993. *Lambda calculi with types*. Oxford University Press, Inc., USA, 117–309.
- [6] Nick Benton, Chung-Kil Hur, Andrew J. Kennedy, and Conor McBride. 2012. Strongly Typed Term Representations in Coq. *J. Autom. Reason.* 49, 2 (Aug. 2012), 141–159. doi:10.1007/s10817-011-9219-0
- [7] Jean-Philippe Bernardy and Nicolas Pouillard. 2013. Names for free: polymorphic views of names and binders. In *Proceedings of the 2013 ACM SIGPLAN Symposium on Haskell, Boston, MA, USA, September 23-24, 2013*, Chung chieh Shan (Ed.). ACM, 13–24. doi:10.1145/2503778.2503780
- [8] Richard S. Bird and Ross Paterson. 1999. de Bruijn notation as a nested datatype. *Journal of Functional Programming* 9, 1 (1999), 77–91. doi:10.1017/S0956796899003366
- [9] Manuel M. T. Chakravarty, Gabriele Keller, Simon Peyton Jones, and Simon Marlow. 2005. Associated types with class. *SIGPLAN Not.* 40, 1 (Jan. 2005), 1–13. doi:10.1145/1047659.1040306

- [10] Adam Chlipala. 2008. Parametric higher-order abstract syntax for mechanized semantics. In *Proceedings of the 13th ACM SIGPLAN International Conference on Functional Programming* (Victoria, BC, Canada) (ICFP '08). Association for Computing Machinery, New York, NY, USA, 143–156. doi:10.1145/1411204.1411226
- [11] David Thrane Christiansen, Iavor S. Diatchki, Robert Dockins, Joe Hendrix, and Tristan Ravitch. 2019. Dependently typed Haskell in industry (experience report). *Proc. ACM Program. Lang.* 3, ICFP, Article 100 (July 2019), 16 pages. doi:10.1145/3341704
- [12] Koen Claessen and John Hughes. 2000. QuickCheck: a lightweight tool for random testing of Haskell programs. In *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming* (ICFP '00). Association for Computing Machinery, New York, NY, USA, 268–279. doi:10.1145/351240.351266
- [13] N.G de Bruijn. 1972. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. *Indagationes Mathematicae (Proceedings)* 75, 5 (1972), 381–392. doi:10.1016/1385-7258(72)90034-0
- [14] Richard A. Eisenberg. 2020. Stitch: the sound type-indexed type checker (functional pearl). In *Proceedings of the 13th ACM SIGPLAN International Symposium on Haskell* (Virtual Event, USA) (Haskell 2020). Association for Computing Machinery, New York, NY, USA, 39–53. doi:10.1145/3406088.3409015
- [15] Richard A. Eisenberg and Stephanie Weirich. 2012. Dependently typed programming with singletons. In *Haskell Symposium*. Copenhagen, Denmark, 117–130.
- [16] Gergő Érdi. 2018. Generic Description of Well-Scoped, Well-Typed Syntaxes. arXiv:1804.00119 [cs.PL] <https://arxiv.org/abs/1804.00119>
- [17] Jean-Yves Girard. 1972. *Interprétation fonctionnelle et élimination des coupures de l'arithmétique d'ordre supérieur*. Thèse d'état. Université de Paris 7.
- [18] Louis-Julien Guillemette and Stefan Monnier. 2008. A type-preserving compiler in Haskell. *SIGPLAN Not.* 43, 9 (Sept. 2008), 75–86. doi:10.1145/1411203.1411218
- [19] Mark P. Jones. 1995. Functional Programming with Overloading and Higher-Order Polymorphism. In *First International Spring School on Advanced Functional Programming Techniques, Båstad, Sweden*. Number 925 in Lecture Notes in Computer Science. Springer-Verlag.
- [20] Aleksey Kliger and Austin Erlandson. 2014. *unbound-generics: Support for programming with names and binders using GHC Generics*. <https://github.com/lambda4geek/unbound-generics>
- [21] Edward A. Kmett. 2013. The bound package. <https://github.com/ekmett/bound/>
- [22] Nikolai Kudasov, Renata Shakirova, Egor Shalagin, and Karina Tyulebaeva. 2024. Free Foil: Generating Efficient and Scope-Safe Abstract Syntax. In *2024 4th International Conference on Code Quality (ICCCQ)*. 1–16. doi:10.1109/ICCCQ60895.2024.10576867
- [23] Peter J. Landin. 1964. The Mechanical Evaluation of Expressions. *Computer Journal* 6, 4 (Jan. 1964), 308–320.
- [24] Andres Löb, Conor McBride, and Wouter Swierstra. 2010. A Tutorial Implementation of a Dependently Typed Lambda Calculus. *Fundam. Inf.* 102, 2 (April 2010), 177–207.
- [25] Dougal Maclaurin, Alexey Radul, and Adam Paszke. 2023. The Foil: Capture-Avoiding Substitution With No Sharp Edges. In *Proceedings of the 34th Symposium on Implementation and Application of Functional Languages* (Copenhagen, Denmark) (IFL '22). Association for Computing Machinery, New York, NY, USA, Article 8, 10 pages. doi:10.1145/3587216.3587224
- [26] José Pedro Magalhães, Atze Dijkstra, Johan Jeuring, and Andres Löb. 2010. A generic deriving mechanism for Haskell. In *Proceedings of the Third ACM Haskell Symposium on Haskell* (Baltimore, Maryland, USA) (Haskell '10). Association for Computing Machinery, New York, NY, USA, 37–48. doi:10.1145/1863523.1863529

- [27] Kazutaka Matsuda, Samantha Frohlich, Meng Wang, and Nicolas Wu. 2023. Embedding by Unembedding. *Proc. ACM Program. Lang.* 7, ICFP, Article 189 (Aug. 2023), 47 pages. doi:10.1145/3607830
- [28] Conor McBride. 1999. *Dependently Typed Functional Programs and their Proofs*. Ph. D. Dissertation. University of Edinburgh.
- [29] Conor McBride and James McKinna. 2004. Functional pearl: I am not a number—I am a free variable. In *Proceedings of the 2004 ACM SIGPLAN Workshop on Haskell* (Snowbird, Utah, USA) (*Haskell '04*). Association for Computing Machinery, New York, NY, USA, 1–9. doi:10.1145/1017472.1017477
- [30] Simon Peyton Jones and Simon Marlow. 2002. Secrets of the Glasgow Haskell Compiler inliner. *J. Funct. Program.* 12, 5 (July 2002), 393–434. doi:10.1017/S0956796802004331
- [31] Steven Schäfer, Tobias Tebbi, and Gert Smolka. 2015. Autosubst: Reasoning with de Bruijn Terms and Parallel Substitutions. In *Interactive Theorem Proving*, Christian Urban and Xingyuan Zhang (Eds.). Springer International Publishing, Cham, 359–374.
- [32] Tim Sheard and Simon Peyton Jones. 2002. Template meta-programming for Haskell. In *Proceedings of the 2002 ACM SIGPLAN Workshop on Haskell* (Pittsburgh, Pennsylvania) (*Haskell '02*). Association for Computing Machinery, New York, NY, USA, 1–16. doi:10.1145/581690.581691
- [33] Kathrin Stark, Steven Schäfer, and Jonas Kaiser. 2019. Autosubst 2: reasoning with multi-sorted de Bruijn terms and vector substitutions. In *Proceedings of the 8th ACM SIGPLAN International Conference on Certified Programs and Proofs* (Cascais, Portugal) (*CPP 2019*). Association for Computing Machinery, New York, NY, USA, 166–180. doi:10.1145/3293880.3294101
- [34] Philip Wadler. 2024. Explicit Weakening. In *A Second Soul: Celebrating the Many Languages of Programming - Festschrift in Honor of Peter Thiemann's Sixtieth Birthday Freiburg, Germany, 30th August 2024*. Number 413. Electronic Proceedings in Theoretical Computer Science, 15–26. doi:10.4204/EPTCS.413
- [35] Geoffrey Washburn and Stephanie Weirich. 2008. Boxes Go Bananas: Encoding Higher-order Abstract Syntax with Parametric Polymorphism. *Journal of Functional Programming* 18, 1 (Jan. 2008), 87–140.
- [36] Stephanie Weirich. 2023. Implementing Dependent Types in pi-forall. doi:10.48550/ARXIV.2207.02129 Lecture notes for the Oregon Programming Languages Summer School.
- [37] Stephanie Weirich, Brent A. Yorgey, and Tim Sheard. 2011. Binders Unbound. In *Proceeding of the 16th ACM SIGPLAN International Conference on Functional Programming* (Tokyo, Japan) (*ICFP '11*). New York, NY, USA, 333–345.